ArXiv Research Field Activity Analysis Using Hadoop MapReduce

University of Ruhuna — Faculty of Engineering
EE7222/EC7204 Cloud Computing | Group 46
Ashfaq MRM (EG/2021/4417) Dissanayake WMNP (EG/2021/4499) Munsif MFA (EG/2021/4684)

March 15, 2026

1. Objective

This project analyzes the ArXiv metadata snapshot using Hadoop MapReduce in a pseudo-distributed single-node environment. For each research category the job computes total paper count, average version count, average abstract length, and average author count. The dataset is naturally suited for MapReduce because each paper can belong to multiple categories, forming a many-to-many aggregation problem across 2.2 million records.

2. Dataset and Environment

The input is the Cornell University ArXiv metadata snapshot (Kaggle), stored in HDFS as JSON Lines — one complete JSON object per line. Each record contains: `categories`, `versions`, `abstract`, and `authors_parsed`. The implementation ran in WSL2 (Ubuntu 22.04) using Java 11, Hadoop 3.3.6, HDFS, YARN, and Maven with a fat JAR built via the assembly plugin.

3. MapReduce Design

Pipeline Overview:

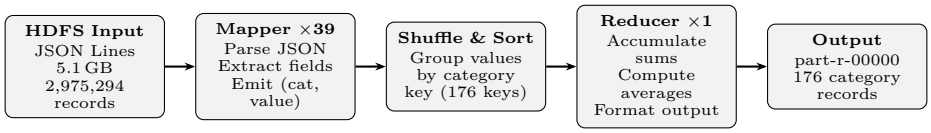

Mapper parses one JSON object per line, extracts category codes, version count, abstract word count, and author count. For every category a paper belongs to, it emits:

`(category, "1 | versionCount | abstractWords | authorCount")`

This allows one paper to contribute to multiple fields (e.g. `cs.LG` and `cs.AI`) simultaneously.

Reducer aggregates raw counts per category and computes:

$$\text{avgVersions} = \frac{\sum v_i}{N}, \quad \text{avgAbstractWords} = \frac{\sum w_i}{N}, \quad \text{avgAuthors} = \frac{\sum a_i}{N}$$

where N is the total paper count for that category. Averages are deferred entirely to the reducer, making the design combiner-safe with no intermediate precision loss.

4. Execution Summary

The final YARN run completed in approximately 2 minutes and 11 seconds. Key job counters:

- Map input records: 2,975,294 | Map tasks: 39 | Reduce tasks: 1
- HDFS bytes read: 5,142,915,425 | Throughput: ≈22,700 records/sec
- Total categories in output: 176 | Malformed/skipped records: 0

5. Interpretation of Results

Patterns & Insights. The output reveals four behavioral axes across 176 categories. (i) *Scale:* `cs.LG` (256K papers) is the largest category, driven by the post-2015 deep learning boom, followed by `hep-ph` (194K) and

Table 1: Representative Output — Selected to Highlight Metric Extremes Across Fields

Category	Total Papers	Avg Versions	Avg Words	Avg Authors
`cs.LG`	256,056	1.70	172	4.4
`cs.CV`	183,697	1.60	177	5.1
`hep-ph`	194,235	1.71	124	3.7
`hep-th`	180,439	1.86	121	2.4
`hep-ex`	59,386	1.70	130	**33.5**
`astro-ph.GA`	75,811	1.42	**217**	9.1
`math.AG`	58,593	1.93	**92**	1.8
`econ.EM`	5,149	**2.12**	152	2.5

`cs.CV` (184K). (ii) *Collaboration:* The author-count axis shows the widest spread — experimental physics (`hep-ex`: 33.5 authors, `nucl-ex`: 16.9, `astro-ph.IM`: 14.0) reflects large-detector consortia (LHC, ALMA), while pure mathematics (`math.AG`, `math.DG`) averages 1.8–1.9, typical of solo theoretical work. (iii) *Abstract length:* Astrophysics sub-categories write the longest abstracts (`astro-ph.GA`: 217 words, `astro-ph.EP`: 209) because observational papers describe instruments and datasets, whereas pure math abstracts (`math.AG`: 92, `math.CV`: 89) rely on symbolic notation over prose. (iv) *Revision behavior:* Economics has the highest revision rates (`econ.EM`: 2.12, `econ.TH`: 2.03), reflecting iterative peer-review cycles, while legacy deprecated categories (`acc-phys`: 49 papers, avg 1.10; `plasm-ph`: 38 papers, avg 1.13) show near-zero revision activity.

Performance. On a single WSL2 node the job sustained ≈22,700 records/sec and read 5.1 GB in 131 seconds. The 39 map splits arose automatically from Hadoop's default 128 MB block size. The single reducer handled 176 unique keys and completed in under 10 seconds of reduce time. Error counters recorded zero malformed JSON lines and zero missing-category records, confirming full dataset coverage.

6. Future Extensions

- **Temporal trends:** Partition mapper output by submission year (from `versions[0].created`) to compute year-over-year growth rates per category.
- **Distributed scale-out:** Deploy on a multi-node YARN cluster (AWS EMR or GCP Dataproc) with multiple reducers partitioned by category prefix to remove the single-reducer bottleneck.
- **NLP enrichment:** Chain a second MapReduce job on the `abstract` field to compute per-category TF-IDF keyword scores or latent semantic topics, enabling automatic field labeling.

7. MapReduce Logic Pseudocode

```
— MAPPER —  map(key, jsonLine):
  if line empty or JSON invalid ⇒ increment errorCounter; return
  cats ← json["categories"].split(" ");  if empty ⇒ skip
  v ← versions.length;  w ← wordCount(abstract);  a ← authors_parsed.length
  for each cat in cats:  emit(cat, "1|v|w|a")
— REDUCER —  reduce(category, values):
  N, sumV, sumW, sumA ← 0
  for each val in values:  parse (c,v,w,a); N+=c;  sumV+=v;  sumW+=w;  sumA+=a
  if N=0:  return;  emit(category, "total=N | avg_ver=sumV/N | avg_words=sumW/N | avg_authors=sumA/N")
```